

Parallel Automata Processing Review

Jose Chavez

Department of Computer Science
University of Engineering and Technology

Operating Systems, 2019

Main
Concepts

Finite State
Machines (FSM)
The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution
Parallel FSM

Overall
Speedup

① Main Concepts

Finite State Machines (FSM)
The Micron Automata Processor (AP)

② Problems

③ Solution

Theoretical Solution
Parallel FSM

④ Overall Speedup

Finite State Machines I

FSM definition

- Formally described by a quintuple $\langle Q, \Sigma, \delta, q_0, F \rangle$.
- Q is a set of states.
- Σ is the input symbol alphabet.
- $\delta(Q, \alpha)$ is the transition function
- q_0 is the set of start states.
- F is the set of reporting or accepting states.




Figure: State representation(NFA)[2]

Finite State Machines II

FSM definition

Main
Concepts

Finite State
Machines (FSM)

The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution

Parallel FSM

Overall
Speedup

T:	/	*	x
a	b	a	a
b	b	c	a
c	c	d	c
d	a	d	c

Figure: Transition Table representation[2]

The Micron Automata Processor (AP) I

Main
Concepts

Finite State
Machines (FSM)

The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution

Parallel FSM

Overall
Speedup



The Micron Automata Processor (AP) II

Main Concepts

Finite State
Machines (FSM)

The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution

Parallel FSM

Overall Speedup

- "Microns 48-chip evaluation board scales this bandwidth to a ridiculous 38TB/s, which enables Automata to solve problems that traditional processors cannot." - *Micron*



Problems

Main Concepts

Finite State
Machines (FSM)
The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution
Parallel FSM

Overall Speedup

- These embarrassingly sequential applications (FSM) with irregular memory access patterns perform poorly on conventional von-Neumann architectures.
- FSM computation, especially Non-Deterministic Finite Automata (NFA) computation is inherently hard to speedup.
- Modern multi-core processors are limited by the number of transitions they can do per thread in a given cycle, limiting the number of patterns they can identify.

Theoretical Solution I

Partitioning the input string into segments and processing these segments concurrently.

The problem with this approach is that starting states for each segment are unknown except the first segment.

- 1 Represent the NFA with a compact equal form, that is the AutomataNetwork Markup Language (ANML) NFA
- 2 We look up for destination states by the transition table and we convert them into active states for the next step.
- 3 We establish the routing matrix and use it to store the transitions and the function.
- 4 Determine the set of states that are active in a particular cycle.

Parallel FSM I

Main Concepts

Finite State
Machines (FSM)

The Micron
Automata Processor
(AP)

Problems

Solution

Theoretical Solution

Parallel FSM

Overall Speedup

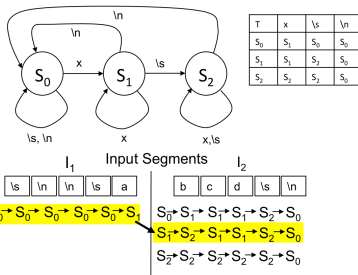


Figure: A FSM example with enumeration[1]

Base Enumerative Technique

Stores a list of states at each step, insted of a single state.

Overall Speedup I

With these implementations we reach a **theoretical** **$2 \times (\text{number of partitions})$** over sequential baseline.
And in a **experimental** result we reached a significant speedup of **25.5x** on average compared again with sequential execution.

Reference I



Aron Subramaniyan

Parallel Automata Processor.

ISCA '17 June.



Todd Mytkowicz.

Data-Parallel Finite-State Machines

ASPLOS'14 March.